

# Gradient Descent for Linear Regression

You will:

- Update gradient descent for logistic regression.
- See gradient descent on a familiar data set.

## Logistic Gradient Descent

Recall the gradient descent algorithm utilizes the gradient calculation:

repeat until convergence: {  
     $w_j = w_j - \alpha \frac{\partial J(\mathbf{w}, b)}{\partial w_j}$       for j := 0..n-1  
     $b = b - \alpha \frac{\partial J(\mathbf{w}, b)}{\partial b}$   
}

(1)

Where each iteration performs simultaneous updates on  $w_j$  for all  $j$ , where

$$\frac{\partial J(\mathbf{w}, b)}{\partial w_j} = \frac{1}{m} \sum_{i=0}^{m-1} (f_{\mathbf{w}, b}(\mathbf{x}^{(i)}) - y^{(i)}) x_j^{(i)}$$

(2)

$$\frac{\partial J(\mathbf{w}, b)}{\partial b} = \frac{1}{m} \sum_{i=0}^{m-1} (f_{\mathbf{w}, b}(\mathbf{x}^{(i)}) - y^{(i)})$$

(3)

- m is the number of training examples in the data set
- $f_{\mathbf{w}, b}(x^{(i)})$  is the model's prediction, while  $y^{(i)}$  is the target
- For a logistic regression model  
     $z = \mathbf{w} \cdot \mathbf{x} + b$   
     $f_{\mathbf{w}, b}(x) = g(z)$   
    where  $g(z)$  is the sigmoid function:  
     $g(z) = \frac{1}{1+e^{-z}}$

## Calculating the Gradient, Code Description

Implements equation (2),(3) above for all  $w_j$  and  $b$ . There are many ways to implement this. Outlined below is this:

- initialize variables to accumulate `dj_dw` and `dj_db`
- for each example
  - calculate the error for that example  $g(\mathbf{w} \cdot \mathbf{x}^{(i)} + b) - \mathbf{y}^{(i)}$
  - for each input value  $x_j^{(i)}$  in this example,
    - multiply the error by the input  $x_j^{(i)}$ , and add to the corresponding element of `dj_dw`. (equation 2 above)
  - add the error to `dj_db` (equation 3 above)
- divide `dj_db` and `dj_dw` by total number of examples (m)
- note that  $\mathbf{x}^{(i)}$  in numpy `X[i,:]` or `X[i]` and  $x_j^{(i)}$  is `X[i,j]`

### Python - imports

```
import numpy as np
import matplotlib.pyplot as plt
import math, copy
```

0.5s

Python

### Python - imports

```
def sigmoid(z):
    """
    compute sigmoid

    Args:
    | z (scalar): logistic function, f_wb

    Return:
    | g (scalar): sigmoid of f_wb
    """
    return 1/(1 + np.exp(-z))
```

0.0s

Python

### C - imports

```
1 double sigmoid(double z){
2     /*
3     compute sigmoid
4
5     Args:
6         z (scalar): logistic function, f_wb
7
8     Return:
9         g (scalar): sigmoid of f_wb
10    */
11
12    return 1 / (1 + exp(-z));
13 }
14
```

Python - gradient

```
def compute_gradient_logistic(X, y, w, b):
    """
    Computes the gradient for linear regression

    Args:
        X (ndarray (m,n)): Data, m examples with n features
        y (ndarray (m, )): target values
        w (ndarray (n, )): model parameters
        b (scalar)       : model parameter

    Returns:
        dj_dw (ndarray (n, )): The gradient of the cost w.r.t the parameters w.
        dj_db (scalar)       : The gradient of the cost w.r.t the parameter b.
    """

    m,n = X.shape
    dj_dw = np.zeros((n,))          #(n,)
    dj_db = 0.

    for i in range(m):
        f_wb_i = sigmoid(np.dot(X[i], w) + b)          #(m,n)(n,)=scalar
        err_i = f_wb_i - y[i]                          #scalar
        for j in range(n):
            dj_dw[j] = dj_dw[j] + err_i * X[i,j]        #scalar
        dj_db = dj_db + err_i
    dj_dw = dj_dw/m                                #(n,)
    dj_db = dj_db/m                                #scalar

    return dj_dw,dj_db
```

✓ 0.0s Python

C - gradient

```
1 struct caches gradient_descent(float x_in[][n], int y_in[], double **w_in, double b_in, float alpha, int num_iters){
2
3     /*
4     Performs batch gradient descent
5
6     Args:
7     X (ndarray (m,n) : Data, m examples with n features
8     y (ndarray (m,)) : target values
9     w_in (ndarray (n,)): Initial values of model parameters
10    b_in (scalar) : Initial values of model parameter
11    alpha (float) : Learning rate
12    num_iters (scalar) : number of iterations to run gradient descent
13
14    Returns:
15    w (ndarray (n,)) : Updated values of parameters
16    b (scalar) : Updated value of parameter
17    */
18
19    struct caches cache;
20    struct Grads grads;
21
22    cache.J_history = (double *)calloc(num_iters, sizeof(double));
23    cache.w = deepcopy(w_in); //avoid modifying global w within function
24    cache.b=b_in;
25
26    for (int i = 0; i < num_iters; i++)
27    {
28        //calculate gradient and update the parameters
29        grads=compute_gradient_logistic(x_in, y_in, cache.w, cache.b);
30
31        //update parameters using w, b, alpha and gradient
32        for (int j = 0; j < n; j++)
33        {
34            cache.w[j][0] = cache.w[j][0] - (grads.djdw[j] * alpha);
35        }
36        cache.b = cache.b - (grads.djdb * alpha);
37
38        double J=compute_cost_logistic(x_in, y_in, cache.w, cache.b);
39
40        cache.J_history[i] = J;
41
42        if (i % (num_iters / 10)==0)
43        {
44            printf("Iteration %d \t Cost %f\n", i, cache.J_history[i]);
45        }
46    }
47
48    return cache;
49 }
50 }
```

Python - cost

```
def compute_cost_logistic(X, y, w, b):
    """
    Compute cost
    Args:
        X (ndarray (m,n)): input data
        y (ndarray (m,)): output data
        w (ndarray (n,)): model parametrs
        b (scalar): model parameter

    Return:
        J (sclar): Cost of logistic model
    """

    J = 0.0
    f_wb_i = 0.0
    m = X.shape[0]

    for i in range(m):
        f_wb_i = sigmoid(np.dot(X[i], w) + b)          #(m,): scalar
        J += -y[i]*np.log(f_wb_i) - (1 - y[i])*np.log(1 - f_wb_i)
    J /= m

    return J
```

✓ 0.0s Python

C - cost

```
1 double compute_cost_logistic(float x[][n], int y[], double **w, double b){
2
3     /*
4     Compute cost
5     Args:
6     X (ndarray (m,n)): input data
7     y (ndarray (m,)): output data
8     w (ndarray (n,)): model parametrs
9     b (scalar): model parameter
10
11    Return:
12    J (sclar): Cost of logistic model
13    */
14
15    double J=0;
16    double f_x;
17    double f_wb;
18
19    for (int i = 0; i < m; i++)
20    {
21        for (int j = 0; j < 1; j++)
22        {
23            f_wb = 0;
24            f_x=0;
25            for (int k = 0; k < n; k++)
26            {
27                f_x += x[i][k] * w[k][j];
28            }
29            f_x +=b;
30            f_wb = sigmoid(f_x);
31            J += (-y[i]*log(f_wb)) - ((1 - y[i])*log(1 - f_wb));
32        }
33    }
34
35    }
36    J /= m;
37
38    return J;
39 }
40 }
```

```
def gradient_descent(X, y, w_in, b_in, alpha, num_iter):
    """
    Performs batch gradient descent

    Args:
    X (ndarray (m,n)) : Data, m examples with n features
    y (ndarray (m,)) : target values
    w_in (ndarray (n,)): Initial values of model parameters
    b_in (scalar) : Initial values of model parameter
    alpha (float) : Learning rate
    num_iters (scalar) : number of iterations to run gradient descent

    Returns:
    w (ndarray (n,)) : Updated values of parameters
    b (scalar) : Updated value of parameter
    """

    #An array to store cost J and w's each iteration primarily for graphing later
    J_history = []
    w = copy.deepcopy(w_in) #avoid modifying global w within function
    b = b_in

    for i in range(num_iter):
        #calculate gradient and update the parameters
        dj_dw, dj_db = compute_gradient_logistic(X, y, w, b)

        #Update Parameters using w, b, alpha and gardient
        w = w - alpha * dj_dw
        b = b - alpha * dj_db

        #Save cost J each iteration
        if i<100000: #prevent resource exhaustion
            J_history.append(compute_cost_logistic(X, y, w, b))

        #Print cost every at intervals 10 times or as many iterations if < 10
        if i% math.ceil(num_iter / 10) ==0:
            print(f"Iteration {i:4d} Cost {J_history[-1]} ")

    return w, b, J_history #return final w, b and J history for graphing
```

```
1 struct caches gradient_descent(float x_in[][n], int y_in[], double **w_in, double b_in, float alpha, int num_iters){
2
3     /*
4     Performs batch gradient descent
5
6     Args:
7     X (ndarray (m,n)) : Data, m examples with n features
8     y (ndarray (m,)) : target values
9     w_in (ndarray (n,)): Initial values of model parameters
10    b_in (scalar) : Initial values of model parameter
11    alpha (float) : Learning rate
12    num_iters (scalar) : number of iterations to run gradient descent
13
14    Returns:
15    w (ndarray (n,)) : Updated values of parameters
16    b (scalar) : Updated value of parameter
17    */
18
19    struct caches cache;
20    struct Grads grads;
21
22    cache.J_history = (double *)calloc(num_iters, sizeof(double));
23    cache.w = deepcopy(w_in); //avoid modifying global w within function
24    cache.b=b_in;
25
26    for (int i = 0; i < num_iters; i++)
27    {
28        //calculate gradient and update the parameters
29        grads=compute_gradient_logistic(x_in, y_in, cache.w, cache.b);
30
31        //update parameters using w, b, alpha and gradient
32        for (int j = 0; j < n; j++)
33        {
34            cache.w[j][0] = cache.w[j][0] - (grads.djdw[j] * alpha);
35        }
36        cache.b = cache.b - (grads.djdb * alpha);
37
38        double J=compute_cost_logistic(x_in, y_in, cache.w, cache.b);
39
40        cache.J_history[i] = J;
41
42        if (i % (num_iters / 10)==0)
43        {
44            printf("Iteration %d \t Cost %f\n", i, cache.J_history[i]);
45        }
46    }
47
48    return cache;
49
50 }
```



Python - main

```
def main():
    X_train = np.array([[0.5, 1.5], [1,1], [1.5, 0.5], [3, 0.5], [2, 2], [1, 2.5]])
    y_train = np.array([0, 0, 0, 1, 1, 1])

    w_tmp = np.zeros_like(X_train[0])
    b_tmp = 0.
    alph = 0.1
    iters = 10000000

    w_out, b_out, _=gradient_descent(X_train, y_train, w_tmp, b_tmp, alph, iters)
    print(f"\nupdated parameters: w:{w_out},    b:{b_out}")

if __name__=="__main__":
    main()
```

✓ 2m 7.0s Python

Iteration 0 Cost 0.684610468560574  
Iteration 10000 Cost 0.01711604647887364  
Iteration 20000 Cost 0.008523403979166467  
Iteration 30000 Cost 0.005672197191107633  
Iteration 40000 Cost 0.004250161053834308  
Iteration 50000 Cost 0.003398230224179212  
Iteration 60000 Cost 0.0028308425601004327  
Iteration 70000 Cost 0.002425848306579758  
Iteration 80000 Cost 0.0021222573122028584  
Iteration 90000 Cost 0.0018862216652143864

updated parameters: w:[8.35313087 8.15226727], b:-22.690605796630248

C - main

```
1 #include<stdio.h>
2 #include<stdlib.h>
3 #include<math.h>
4
5 #define m 6
6 #define n 2
7
8 struct Grads{
9     double *djdW;
10    double dJdb;
11 } ;
12
13 struct caches{
14     double *J_history;
15     double **w;
16     double b;
17 };
18
19 double **deepcopy(double **src);
20 double compute_cost_logistic(float x[][n], int y[], double **w, double b);
21 struct Grads compute_gradient_logistic(float x[][n], int y[], double **w, double b);
22 struct caches gradient_descent(float x_in[][n], int y_in[], double **w_in, double b_in, float alpha, int num_iters);
23
24
25 int main(){
26
27     float X_train[n] = {{0.5, 1.5}, {1, 1}, {1.5, 0.5}, {3, 0.5}, {2, 2}, {1, 2.5}};
28     int y_train[] = {0, 0, 0, 1, 1, 1};
29
30     double **w_init=(double **)calloc(n, sizeof(double *));
31     if (w_init == NULL)
32     {
33         perror("Error in allocating memory");
34     }
35     for (int i = 0; i < n; i++)
36     {
37         w_init[i] = (double *)calloc(1, sizeof(double));
38         if (w_init[i]==NULL)
39         {
40             perror("Error in allocating memory");
41             for (int j = 0; j < i; j++)
42             {
43                 free(w_init[j]);
44             }
45             free(w_init);
46         }
47         w_init[i][0] = 0;
48     }
49
50 }
51
52 double b_init = 0;
53 double alpha=1e-1;
54 int num_iters=10000000;
55
56 struct caches cache;
57 cache = gradient_descent(X_train, y_train, w_init, b_init, alpha, num_iters);
58
59 printf("updated parameters: ");
60 for (int i = 0; i < n; i++)
61 {
62     printf("w[%d]: %f \t", i, cache.w[i][0]);
63 }
64 printf("b: %f\n", cache.b);
65
66 for (int i = 0; i < n; i++)
67 {
68     free(w_init[i]);
69     free(cache.w[i]);
70 }
71 free(cache.w);
72 free(w_init);
73 free(cache.J_history);
74
75 return 0;
76 }
77
```

● venvsuzanodero@suzans-MacBook-Air Gradient\_Descent % gcc gradient\_descent.c  
● venvsuzanodero@suzans-MacBook-Air Gradient\_Descent % ./a.out  
Iteration 0 Cost 0.684610  
Iteration 10000 Cost 0.017116  
Iteration 20000 Cost 0.008523  
Iteration 30000 Cost 0.005672  
Iteration 40000 Cost 0.004250  
Iteration 50000 Cost 0.003398  
Iteration 60000 Cost 0.002831  
Iteration 70000 Cost 0.002426  
Iteration 80000 Cost 0.002122  
Iteration 90000 Cost 0.001886  
updated parameters: w[0]: 8.353131 w[1]: 8.152267 b: -22.690606